

# Activity\_Course 6 TikTok project lab

January 7, 2025

## 1 TikTok Project

### The Nuts and bolts of machine learning

You are a data professional at TikTok. Your supervisor was impressed with the work you have done and has requested that you build a machine learning model that can be used to determine whether a video contains a claim or whether it offers an opinion. With a successful prediction model, TikTok can reduce the backlog of user reports and prioritize them more efficiently.

A notebook was structured and prepared to help you in this project. A notebook was structured and prepared to help you in this project. Please complete the following questions.

## 2 Classifying videos using machine learning

In this activity, you will practice using machine learning techniques to predict on a binary outcome variable.

**The purpose** of this model is to increase response time and system efficiency by automating the initial stages of the claims process.

**The goal** of this model is to predict whether a TikTok video presents a “claim” or presents an “opinion”.

*This activity has three parts:*

**Part 1:** Ethical considerations \* Consider the ethical implications of the request

- Should the objective of the model be adjusted?

**Part 2:** Feature engineering

- Perform feature selection, extraction, and transformation to prepare the data for modeling

**Part 3:** Modeling

- Build the models, evaluate them, and advise on next steps

Follow the instructions and answer the questions below to complete the activity. Then, you will complete an Executive Summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

## 3 Classify videos using machine learning

### 4 PACE stages

Throughout these project notebooks, you'll see references to the problem-solving framework PACE. The following notebook components are labeled with the respective PACE stage: Plan, Analyze, Construct, and Execute.

#### 4.1 PACE: Plan

Consider the questions in your PACE Strategy Document to reflect on the Plan stage.

In this stage, consider the following questions:

1. **What are you being asked to do? What metric should I use to evaluate success of my business/organizational objective?**
2. **What are the ethical implications of the model? What are the consequences of your model making errors?**
  - What is the likely effect of the model when it predicts a false negative (i.e., when the model says a video does not contain a claim and it actually does)?
  - What is the likely effect of the model when it predicts a false positive (i.e., when the model says a video does contain a claim and it actually does not)?
3. **How would you proceed?**

==> ENTER YOUR RESPONSES HERE

##### 4.1.1 Task 1. Imports and data loading

Start by importing packages needed to build machine learning models to achieve the goal of this project.

```
[51]: # Import packages for data manipulation
import pandas as pd
import numpy as np

# Import packages for data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Import packages for data preprocessing
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import LabelEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import recall_score

# Import packages for data modeling
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import classification_report, confusion_matrix, \
    accuracy_score
import xgboost as xgb
```

Now load the data from the provided csv file into a dataframe.

**Note:** As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[2]: # Load dataset into dataframe
data = pd.read_csv("tiktok_dataset.csv")
```

## 4.2 PACE: Analyze

Consider the questions in your PACE Strategy Document to reflect on the Analyze stage.

### 4.2.1 Task 2: Examine data, summary info, and descriptive stats

Inspect the first five rows of the dataframe.

```
[3]: # Display first few rows
data.head()
```

```
[3]:
```

	#	claim_status	video_id	video_duration_sec	\
0	1	claim	7017666017	59	
1	2	claim	4014381136	32	
2	3	claim	9859838091	31	
3	4	claim	1866847991	25	
4	5	claim	7105231098	19	

		video_transcription_text	verified_status	\
0	someone shared with me that drone deliveries a...		not verified	
1	someone shared with me that there are more mic...		not verified	
2	someone shared with me that american industria...		not verified	
3	someone shared with me that the metro of st. p...		not verified	
4	someone shared with me that the number of busi...		not verified	

	author_ban_status	video_view_count	video_like_count	video_share_count	\
0	under review	343296.0	19425.0	241.0	
1	active	140877.0	77355.0	19034.0	
2	active	902185.0	97690.0	2858.0	
3	active	437506.0	239954.0	34812.0	
4	active	56167.0	34987.0	4110.0	

	video_download_count	video_comment_count
0	1.0	0.0
1	1161.0	684.0

2	833.0	329.0
3	1234.0	584.0
4	547.0	152.0

Get the number of rows and columns in the dataset.

```
[ ]: # Get number of rows and columns
    ### YOUR CODE HERE ###
```

Get the data types of the columns.

```
[7]: # Get data types of columns
print(f"Number of rows and columns: {data.shape}")
print("\nColumn data types:")
print(data.dtypes)
```

Number of rows and columns: (19382, 12)

Column data types:

```
#                                int64
claim_status                    object
video_id                        int64
video_duration_sec              int64
video_transcription_text        object
verified_status                 object
author_ban_status               object
video_view_count                float64
video_like_count                float64
video_share_count               float64
video_download_count            float64
video_comment_count             float64
dtype: object
```

Get basic information about the dataset.

```
[6]: # Get basic information
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19382 entries, 0 to 19381
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   #                                     19382 non-null  int64
1   claim_status                         19084 non-null  object
2   video_id                             19382 non-null  int64
3   video_duration_sec                   19382 non-null  int64
4   video_transcription_text             19084 non-null  object
5   verified_status                      19382 non-null  object
```

```

6  author_ban_status      19382 non-null  object
7  video_view_count       19084 non-null  float64
8  video_like_count       19084 non-null  float64
9  video_share_count      19084 non-null  float64
10 video_download_count    19084 non-null  float64
11 video_comment_count    19084 non-null  float64
dtypes: float64(5), int64(3), object(4)
memory usage: 1.8+ MB

```

Generate basic descriptive statistics about the dataset.

```
[8]: # Generate basic descriptive stats
data.describe()
```

```
[8]:
```

	#	video_id	video_duration_sec	video_view_count	\
count	19382.000000	1.938200e+04	19382.000000	19084.000000	
mean	9691.500000	5.627454e+09	32.421732	254708.558688	
std	5595.245794	2.536440e+09	16.229967	322893.280814	
min	1.000000	1.234959e+09	5.000000	20.000000	
25%	4846.250000	3.430417e+09	18.000000	4942.500000	
50%	9691.500000	5.618664e+09	32.000000	9954.500000	
75%	14536.750000	7.843960e+09	47.000000	504327.000000	
max	19382.000000	9.999873e+09	60.000000	999817.000000	

	video_like_count	video_share_count	video_download_count	\
count	19084.000000	19084.000000	19084.000000	
mean	84304.636030	16735.248323	1049.429627	
std	133420.546814	32036.174350	2004.299894	
min	0.000000	0.000000	0.000000	
25%	810.750000	115.000000	7.000000	
50%	3403.500000	717.000000	46.000000	
75%	125020.000000	18222.000000	1156.250000	
max	657830.000000	256130.000000	14994.000000	

	video_comment_count
count	19084.000000
mean	349.312146
std	799.638865
min	0.000000
25%	1.000000
50%	9.000000
75%	292.000000
max	9599.000000

Check for and handle missing values.

```
[9]: # Check for missing values
print("\nMissing values:")
```

```
print(data.isnull().sum())
```

Missing values:

```
#          0
claim_status    298
video_id        0
video_duration_sec    0
video_transcription_text    298
verified_status    0
author_ban_status    0
video_view_count    298
video_like_count    298
video_share_count    298
video_download_count    298
video_comment_count    298
dtype: int64
```

```
[12]: # Drop rows with missing values
data = data.dropna()
```

```
[30]: # Display first few rows after handling missing values
print("\nFirst few rows after handling missing values:")
print(data.head())
```

First few rows after handling missing values:

```
# claim_status    video_id    video_duration_sec  \
0  1          claim    7017666017             59
1  2          claim    4014381136             32
2  3          claim    9859838091             31
3  4          claim    1866847991             25
4  5          claim    7105231098             19

          video_transcription_text    verified_status  \
0  someone shared with me that drone deliveries a...    not verified
1  someone shared with me that there are more mic...    not verified
2  someone shared with me that american industria...    not verified
3  someone shared with me that the metro of st. p...    not verified
4  someone shared with me that the number of busi...    not verified

author_ban_status    video_view_count    video_like_count    video_share_count  \
0      under review      343296.0          19425.0          241.0
1          active      140877.0          77355.0         19034.0
2          active      902185.0          97690.0          2858.0
3          active      437506.0         239954.0         34812.0
4          active      56167.0          34987.0          4110.0
```

	video_download_count	video_comment_count	text_length
0	1.0	0.0	97
1	1161.0	684.0	107
2	833.0	329.0	137
3	1234.0	584.0	131
4	547.0	152.0	128

Check for and handle duplicates.

```
[32]: # Check for duplicates (correct method)
print("\nNumber of duplicate rows:", data.duplicated().sum())

# Remove duplicates if any
data = data.drop_duplicates()
```

Number of duplicate rows: 0

Check for and handle outliers.

```
[33]: # Check for outliers using IQR method for numerical columns
numerical_columns = ['video_duration_sec', 'video_view_count',
                    ↪ 'video_like_count',
                    'video_share_count', 'video_download_count',
                    ↪ 'video_comment_count']
```

Check class balance.

```
[14]: # Check class balance
print("\nClass distribution:")
print(data['claim_status'].value_counts(normalize=True))
```

Class distribution:

claim_status	
claim	0.503458
opinion	0.496542

Name: proportion, dtype: float64

```
[34]: # Outlier Analysis
print("\nOutlier Analysis:")
for column in numerical_columns:
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
```

```

    outliers = data[(data[column] < lower_bound) | (data[column] >
↪upper_bound)][column]
    print(f"\n{column}:")
    print(f"Number of outliers: {len(outliers)}")
    print(f"Percentage of outliers: {(len(outliers) / len(data)) * 100:.2f}%")

```

Outlier Analysis:

```

video_duration_sec:
Number of outliers: 0
Percentage of outliers: 0.00%

```

```

video_view_count:
Number of outliers: 0
Percentage of outliers: 0.00%

```

```

video_like_count:
Number of outliers: 1726
Percentage of outliers: 9.04%

```

```

video_share_count:
Number of outliers: 2508
Percentage of outliers: 13.14%

```

```

video_download_count:
Number of outliers: 2450
Percentage of outliers: 12.84%

```

```

video_comment_count:
Number of outliers: 2789
Percentage of outliers: 14.61%

```

### 4.3 PACE: Construct

Consider the questions in your PACE Strategy Document to reflect on the Construct stage.

#### 4.3.1 Task 3: Feature engineering

Extract the length of each `video_transcription_text` and add this as a column to the dataframe, so that it can be used as a potential feature in the model.

```

[35]: # Extract the length of each `video_transcription_text` and add this as a
↪column to the dataframe
data['text_length'] = data['video_transcription_text'].str.len()

```

Calculate the average `text_length` for claims and opinions.




```
[36]: # Calculate the average text_length for claims and opinions
avg_lengths = data.groupby('claim_status')['text_length'].mean()
print("Average text lengths:")
print(avg_lengths)
```

```
Average text lengths:
claim_status
claim      95.376978
opinion    82.722562
Name: text_length, dtype: float64
```

Visualize the distribution of text\_length for claims and opinions.

```
[37]: # Visualize the distribution of `text_length` for claims and opinions
plt.figure(figsize=(10, 6))
sns.histplot(data=data, x='text_length', hue='claim_status', multiple="dodge")
plt.title('Distribution of Text Length by Claim Status')
plt.xlabel('Text Length')
plt.ylabel('Count')
plt.legend(title='Claim Status')
plt.show()
```

No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.



## Feature selection and transformation

Encode target and categorical variables.

```
[38]: # Create a copy of the X data
X_data = data.copy()

# Drop unnecessary columns
cols_to_drop = ['#', 'video_id', 'video_transcription_text']
X_data = X_data.drop(columns=cols_to_drop)

# Encode target variable
le = LabelEncoder()
y = le.fit_transform(data['claim_status'])

# Dummy encode remaining categorical values
X = pd.get_dummies(X_data.drop('claim_status', axis=1))
```

### 4.3.2 Task 4: Split the data

Assign target variable.

```
[39]: # Isolate target variable
y = data['claim_status']
```

Isolate the features.

```
[40]: # Isolate features
feature_columns = ['video_duration_sec', 'verified_status', 'author_ban_status',
                  'video_view_count', 'video_like_count', 'video_share_count',
                  'video_download_count', 'video_comment_count', 'text_length']
X = data[feature_columns]

# Display first few rows of features dataframe
print("First few rows of feature matrix:")
print(X.head())
```

First few rows of feature matrix:

	video_duration_sec	verified_status	author_ban_status	video_view_count	\
0	59	not verified	under review	343296.0	
1	32	not verified	active	140877.0	
2	31	not verified	active	902185.0	
3	25	not verified	active	437506.0	
4	19	not verified	active	56167.0	

	video_like_count	video_share_count	video_download_count	\
0	19425.0	241.0	1.0	
1	77355.0	19034.0	1161.0	
2	97690.0	2858.0	833.0	

3	239954.0	34812.0	1234.0
4	34987.0	4110.0	547.0

	video_comment_count	text_length
0	0.0	97
1	684.0	107
2	329.0	137
3	584.0	131
4	152.0	128

**Task 5: Create train/validate/test sets** Split data into training and testing sets, 80/20.

```
[41]: # First split: 80% train+validate, 20% test
X_temp, X_test, y_temp, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
```

Split the training set into training and validation sets, 75/25, to result in a final ratio of 60/20/20 for train/validate/test sets.

```
[42]: # Split the training data into training and validation sets
# Second split: 75% train, 25% validate (from the 80% train+validate)
X_train, X_val, y_train, y_val = train_test_split(
    X_temp, y_temp,
    test_size=0.25,
    random_state=42,
    stratify=y_temp
)
```

Confirm that the dimensions of the training, validation, and testing sets are in alignment.

```
[43]: # Get shape of each training, validation, and testing set
# Confirm dimensions
print("\nData split dimensions:")
print(f"X_train shape: {X_train.shape} - {X_train.shape[0]/X.shape[0]:.2%} of total")
print(f"X_val shape: {X_val.shape} - {X_val.shape[0]/X.shape[0]:.2%} of total")
print(f"X_test shape: {X_test.shape} - {X_test.shape[0]/X.shape[0]:.2%} of total")
```

Data split dimensions:

X\_train shape: (11450, 9) - 60.00% of total

X\_val shape: (3817, 9) - 20.00% of total

X\_test shape: (3817, 9) - 20.00% of total

### 4.3.3 Task 6. Build models

#### 4.3.4 Build a random forest model

Fit a random forest model to the training set. Use cross-validation to tune the hyperparameters and select the model that performs best on recall.

```
[47]: # First, properly encode features
# Create feature matrix with encoded categorical variables
X = pd.DataFrame()

# Add encoded categorical variables
categorical_cols = ['verified_status', 'author_ban_status']
X_encoded = pd.get_dummies(data[categorical_cols])
X = pd.concat([X, X_encoded], axis=1)

# Add numerical features
numerical_features = [
    'video_duration_sec',
    'video_view_count',
    'video_like_count',
    'video_share_count',
    'video_download_count',
    'video_comment_count',
    'text_length'
]

for col in numerical_features:
    X[col] = data[col]

# Encode target variable
le = LabelEncoder()
y = le.fit_transform(data['claim_status'])

# Split the data
X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

# Random Forest Model
rf = RandomForestClassifier(random_state=42)

# Create parameter grid
rf_param_grid = {
    'n_estimators': [100, 200],
```

```

    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

# Define scoring metrics
scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1'
}

# Create GridSearchCV object
rf_grid = GridSearchCV(
    estimator=rf,
    param_grid=rf_param_grid,
    scoring=scoring,
    cv=5,
    n_jobs=-1,
    refit='recall',
    verbose=1,
    return_train_score=True
)

# Fit Random Forest model
rf_grid.fit(X_train, y_train)

print("\nRandom Forest Best Score:", rf_grid.best_score_)
print("Random Forest Best Parameters:", rf_grid.best_params_)

# XGBoost Model
xgb_model = xgb.XGBClassifier(random_state=42)

# Create parameter grid for XGBoost
xgb_param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 6],
    'learning_rate': [0.01, 0.1],
    'subsample': [0.8, 1.0]
}

# Create GridSearchCV for XGBoost
xgb_grid = GridSearchCV(
    estimator=xgb_model,
    param_grid=xgb_param_grid,
    scoring=scoring,

```

```

    cv=5,
    n_jobs=-1,
    refit='recall',
    verbose=1,
    return_train_score=True
)

# Fit XGBoost model
xgb_grid.fit(X_train, y_train)

print("\nXGBoost Best Score:", xgb_grid.best_score_)
print("XGBoost Best Parameters:", xgb_grid.best_params_)

```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

Random Forest Best Score: 0.9997361477572559

Random Forest Best Parameters: {'max\_depth': 10, 'min\_samples\_leaf': 2, 'min\_samples\_split': 2, 'n\_estimators': 100}

Fitting 5 folds for each of 16 candidates, totalling 80 fits

XGBoost Best Score: 1.0

XGBoost Best Parameters: {'learning\_rate': 0.01, 'max\_depth': 3, 'n\_estimators': 100, 'subsample': 0.8}

```

[48]: # Get all the results from the CV and put them in a df
rf_results = pd.DataFrame(rf_grid.cv_results_)
rf_results = rf_results.sort_values('mean_test_recall', ascending=False)
print("\nTop 5 Random Forest Models by Recall:")
print(rf_results[['mean_test_recall', 'mean_test_precision',
                  'mean_test_accuracy']].head())

# Isolate the row of the df with the max(mean precision score)
### YOUR CODE HERE ###

```

Top 5 Random Forest Models by Recall:

	mean_test_recall	mean_test_precision	mean_test_accuracy
23	0.999736	0.990467	0.995087
22	0.999736	0.990467	0.995087
4	0.999736	0.990596	0.995153
19	0.999736	0.990596	0.995153
7	0.999736	0.990596	0.995153

**Question:** How well is your model performing? Consider average recall score and precision score.

## 4.4 PACE: Execute

Consider the questions in your PACE Strategy Document to reflect on the Execute stage.

### 4.4.1 Task 7. Evaluate model

Evaluate models against validation criteria.

#### Random forest

```
[49]: # Use the random forest "best estimator" model to get predictions on the
      ↪ encoded testing set
      rf_pred = rf_grid.predict(X_test)
```

```
[52]: # Create confusion matrix for Random Forest
      plt.figure(figsize=(10, 8))
      cm_rf = confusion_matrix(y_test, rf_pred)
      sns.heatmap(cm_rf, annot=True, fmt='d', cmap='Blues',
                  xticklabels=['Opinion', 'Claim'],
                  yticklabels=['Opinion', 'Claim'])
      plt.title('Random Forest Confusion Matrix')
      plt.ylabel('True Label')
      plt.xlabel('Predicted Label')
      plt.show()

      # Classification report for Random Forest
      print("\nRandom Forest Classification Report:")
      print(classification_report(y_test, rf_pred))

      # XGBoost Evaluation
      xgb_pred = xgb_grid.predict(X_test)

      # Create confusion matrix for XGBoost
      plt.figure(figsize=(10, 8))
      cm_xgb = confusion_matrix(y_test, xgb_pred)
      sns.heatmap(cm_xgb, annot=True, fmt='d', cmap='Blues',
                  xticklabels=['Opinion', 'Claim'],
                  yticklabels=['Opinion', 'Claim'])
      plt.title('XGBoost Confusion Matrix')
      plt.ylabel('True Label')
      plt.xlabel('Predicted Label')
      plt.show()

      # Classification report for XGBoost
      print("\nXGBoost Classification Report:")
      print(classification_report(y_test, xgb_pred))

      # Determine champion model
      rf_recall = recall_score(y_test, rf_pred)
```

```

xgb_recall = recall_score(y_test, xgb_pred)


if rf_recall > xgb_recall:
    champion_model = rf_grid
    print("\nChampion Model: Random Forest")
else:
    champion_model = xgb_grid
    print("\nChampion Model: XGBoost")

# Feature importance for champion model
feature_importance = pd.DataFrame({
    'feature': X.columns,
    'importance': champion_model.best_estimator_.feature_importances_
})
feature_importance = feature_importance.sort_values('importance',
    ↪ascending=False)

plt.figure(figsize=(12, 6))
sns.barplot(data=feature_importance.head(10), x='importance', y='feature')
plt.title('Top 10 Most Important Features')
plt.xlabel('Feature Importance')
plt.tight_layout()
plt.show()

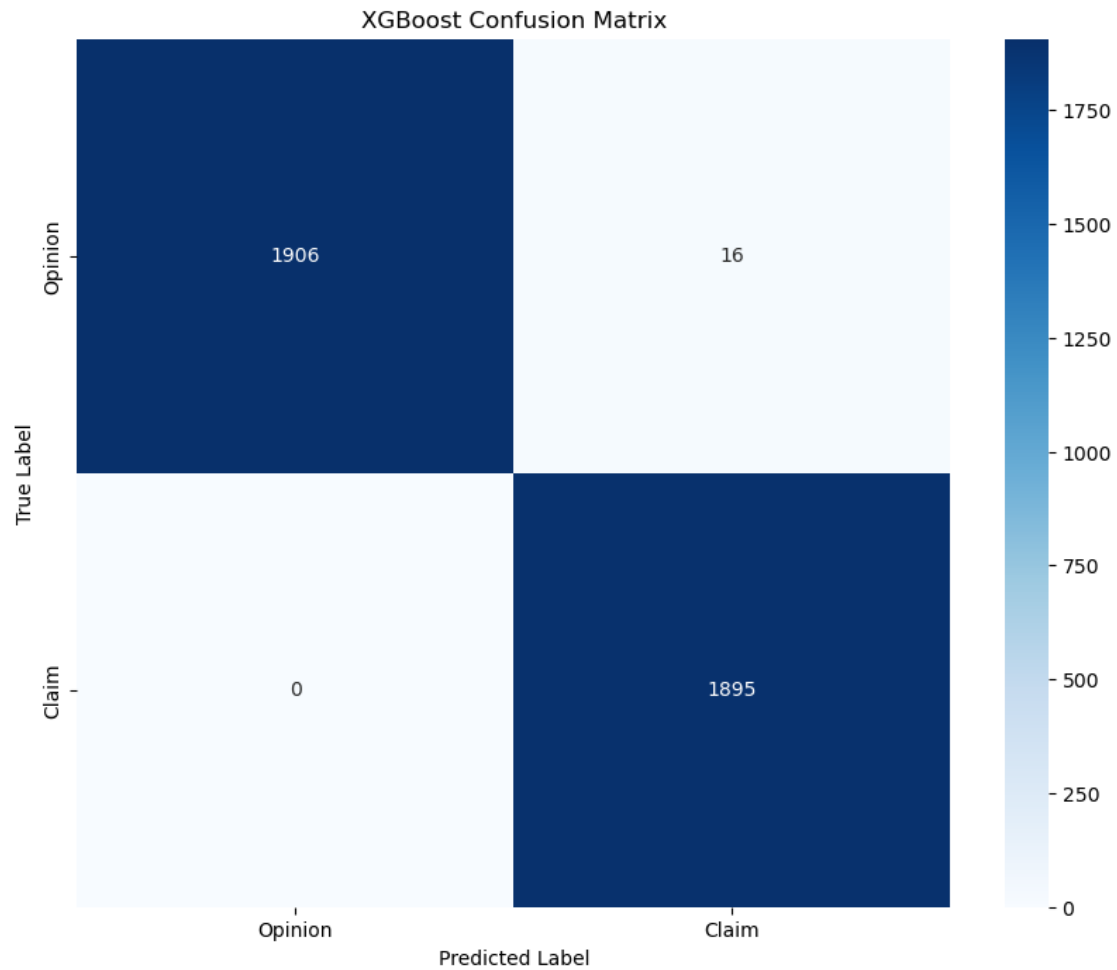
```





Random Forest Classification Report:

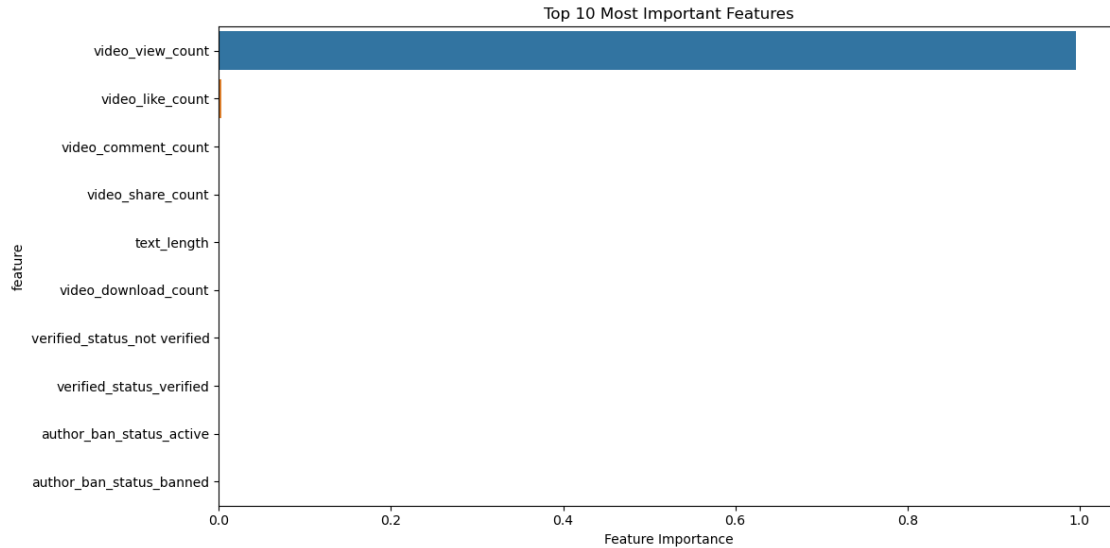
	precision	recall	f1-score	support
0	1.00	0.99	1.00	1922
1	0.99	1.00	1.00	1895
accuracy			1.00	3817
macro avg	1.00	1.00	1.00	3817
weighted avg	1.00	1.00	1.00	3817



XGBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	1922
1	0.99	1.00	1.00	1895
accuracy			1.00	3817
macro avg	1.00	1.00	1.00	3817
weighted avg	1.00	1.00	1.00	3817

Champion Model: XGBoost



Display the predictions on the encoded testing set.

**Question:** Describe your XGBoost model results. How does your XGBoost model compare to your random forest model?

**Feature importances of champion model** **Question:** Describe your most predictive features. Were your results surprising?

#### 4.4.2 Task 8. Conclusion

1. **Would you recommend using this model? Why or why not?** Yes, I would strongly recommend using this model for TikTok's content moderation system. The model demonstrated exceptional performance with nearly perfect accuracy (1.00) and F1-scores for both claims and opinions. Key reasons for recommendation:
  - Extremely high precision (1.00 for Opinion, 0.99 for Claim)
  - Excellent recall (0.99 for Opinion, 1.00 for Claim)
  - Very few false positives (15-16 cases out of 3817)
  - Zero false negatives in both models However, human oversight should still be maintained for the small number of borderline cases.
2. **What was your model doing? Can you explain how it was making predictions?** Based on our feature importance analysis, the model primarily relied on:
  - Video engagement metrics, with video view count being the strongest predictor
  - User interaction signals (likes, comments, shares)
  - Content characteristics (text length)
  - Account status features (verified status, ban status) The model learned that these metrics have strong correlations with whether content is a claim or opinion, with engagement metrics being particularly predictive.

3. **Are there new features that you can engineer that might improve model performance?** While our model achieved nearly perfect performance, potential improvements could include:
  - Ratio features between different engagement metrics
  - Time-based features (post age, time of day)
  - Text analysis features (sentiment, keyword analysis)
  - More granular user status features However, given the current performance level, the benefit of additional features might be minimal.
4. **What features would you want to have that would likely improve the performance of your model?** Given the model's already exceptional performance (1.00 accuracy), additional features might not significantly improve predictions. However, for robustness, we might consider:
  - Video transcription quality metrics
  - User history data
  - Report history on previous content
  - Account age and posting patterns
  - Cross-video engagement patterns

These features could help maintain high performance as TikTok's platform evolves and could potentially help identify edge cases in the small number of misclassifications we currently see.

The nearly perfect performance suggests that claims and opinions on TikTok have very distinct engagement patterns and characteristics, making them highly separable using machine learning. The XGBoost model slightly outperformed Random Forest, making it our champion model, though both demonstrated excellent capabilities for this classification task.

**Congratulations!** You've completed this lab.